ℓ_2^2-learning of vector autoregressive processes

Wessel N. van Wieringen

Dept. of Epidemiology & Data Science, Amsterdam UMC
Dept. of Mathematics, Vrije Universiteit Amsterdam
Amsterdam, the Netherlands

IBC pre-conference course, Riga, 10.07.2022

Auto-regression

The AR(1) process:

$$Y_{t+1} = \alpha Y_t + \varepsilon_{t+1},$$

with $\alpha \in (-1,1)$, $\varepsilon_t \sim_{\text{i.i.d.}} \mathcal{N}(0,\sigma_\varepsilon^2)$ for all t.

Nomenclature:
- α : endogeneous operator, autoregression parameter,
- ε_t : exogeneous signal, innovation.

The *auto-regulator* motif

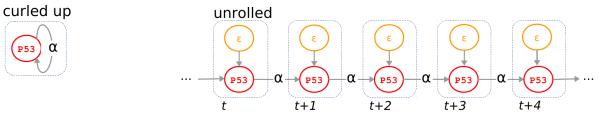

Auto-regression

Stationarity

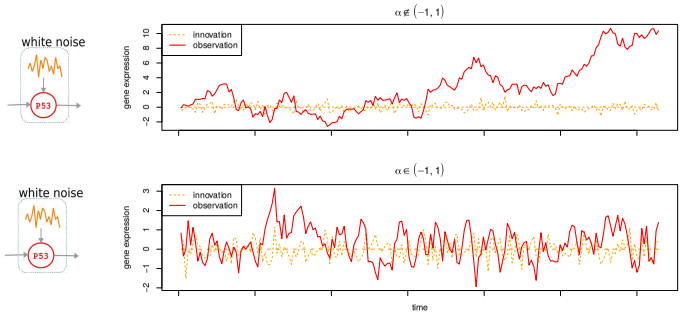

Auto-regression

Pulse processing

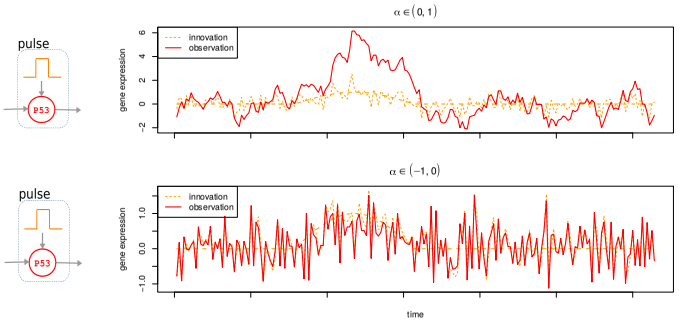

Auto-regression

The AR(1) process can also be formulated as:

$$Y_{t+1} \mid Y_t \sim \mathcal{N}(\alpha Y_t, \sigma_\varepsilon^2).$$

This amounts to a 1st order Markov property:

$$Y_{t+1} \mid Y_t, Y_{t-1}, Y_{t-2}, \ldots \sim Y_{t+1} \mid Y_t.$$

With the 2nd order Markov property, we would arrive at an AR(2) process:

$$Y_{t+1} = \alpha^{(1)} Y_t + \alpha^{(2)} Y_{t-1} + \varepsilon_{t+1},$$

with $\varepsilon_t \sim \mathcal{N}(0,\sigma_\varepsilon^2)$.

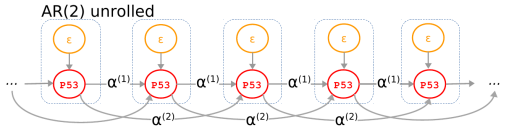

Auto-regression

A VAR(1) process:

$$\begin{aligned} Y_{1,t+1} &= \alpha_{11} Y_{1,t} + \alpha_{12} Y_{2,t} + \varepsilon_{1,t+1}. \\ Y_{2,t+1} &= \alpha_{21} Y_{1,t} + \alpha_{22} Y_{2,t} + \varepsilon_{2,t+1}. \end{aligned}$$

with a stationarity criterion on α's and $(\varepsilon_{1,t},\varepsilon_{2,t})^\top \sim \mathcal{N}(\mathbf{0}_2, \mathbf{\Omega}_\varepsilon^{-1})$.

The *coupled auto-regulators* motif

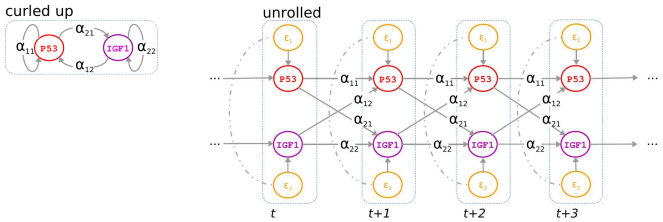

Auto-regression

Coupled auto-regulators may achieve denoising and, necessary for the circadian clock, synchronization.

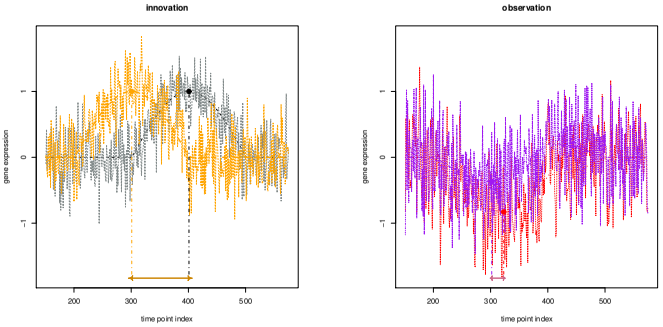

Auto-regression

A VAR(1) process:

$$\begin{aligned} Y_{1,t+1} &= \alpha_{11} Y_{1,t} + \alpha_{12} Y_{2,t} + \alpha_{13} Y_{3,t} + \varepsilon_{1,t+1}, \\ Y_{2,t+1} &= \alpha_{21} Y_{1,t} + \alpha_{22} Y_{2,t} + \alpha_{23} Y_{3,t} + \varepsilon_{2,t+1}, \\ Y_{3,t+1} &= \alpha_{31} Y_{1,t} + \alpha_{32} Y_{2,t} + \alpha_{33} Y_{3,t} + \varepsilon_{3,t+1}. \end{aligned}$$

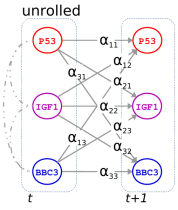

Or, in matrix notation

$$\mathbf{Y}_{t+1} = \mathbf{A}\mathbf{Y}_t + \varepsilon_{t+1}$$

with $|ev_j(\mathbf{A})| < 1$ and $\varepsilon_t \sim \mathcal{N}(\mathbf{0}_3, \mathbf{\Omega}_\varepsilon^{-1})$.

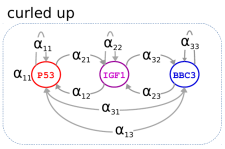

Auto-regression

feedforward loop

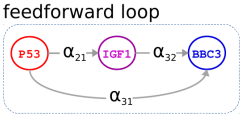

$$\mathbf{A} = \begin{pmatrix} 0 & 0 & 0 \\ \alpha_{21} & 0 & 0 \\ \alpha_{31} & \alpha_{32} & 0 \end{pmatrix}$$

$\mathbf{A} =$ 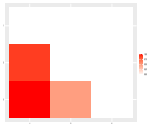

feedback loop

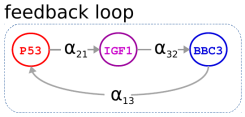

$$\mathbf{A} = \begin{pmatrix} 0 & 0 & \alpha_{13} \\ \alpha_{21} & 0 & 0 \\ 0 & \alpha_{32} & 0 \end{pmatrix}$$

$\mathbf{A} =$ 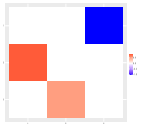

Auto-regression

Pulse processing by the *coherent* feedforward motif

$\mathbf{A} =$ 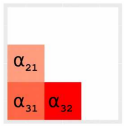

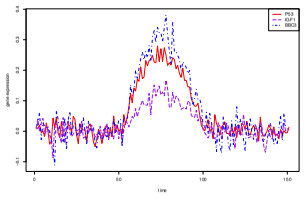

Pulse processing by the *incoherent* feedforward motif

$\mathbf{A} =$ 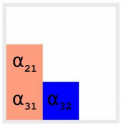

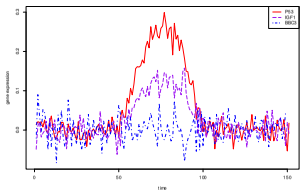

Auto-regression

Systems biology: how will a combination of motifs play out?

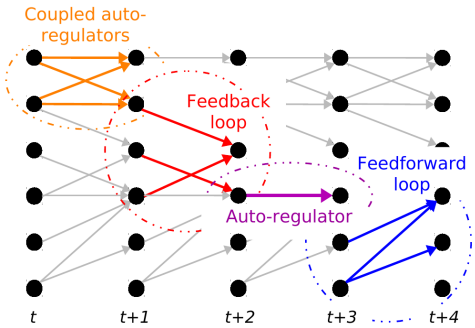

Auto-regression

What can we say about the (marginal) dependence among the variates?

A closer look at the recursive relation:

$$\begin{aligned} \mathbf{Y}_{t+1} &= \mathbf{A}\mathbf{Y}_t + \varepsilon_{t+1} \\ &= \mathbf{A}^2\mathbf{Y}_{t-1} + \mathbf{A}\varepsilon_t + \varepsilon_{t+1} \\ &= \mathbf{A}^3\mathbf{Y}_{t-2} + \mathbf{A}^2\varepsilon_{t-1} + \mathbf{A}\varepsilon_t + \varepsilon_{t+1} \\ &= \dots \end{aligned}$$

Then, e.g.

$$\text{Cov}(\mathbf{Y}_{t+2}, \mathbf{Y}_t) = \text{Cov}(\mathbf{A}^2\mathbf{Y}_t + \mathbf{A}\varepsilon_{t+1} + \varepsilon_{t+2}, \mathbf{Y}_t) = \mathbf{A}^2\mathbf{\Sigma}_y,$$

where $\mathbf{\Sigma}_y$ denotes the process variance and satisfies $\mathbf{\Sigma}_y = \mathbf{A}\mathbf{\Sigma}_y\mathbf{A}^\top + \mathbf{\Sigma}_\varepsilon$.

Nomenclature: *auto-correlation* vs. *cross-correlation*,
$\text{Cov}(Y_{j,t+2}, Y_{j,t})$ vs. $\text{Cov}(Y_{j,t+2}, Y_{j',t})$.

Auto-regression

The *time-series chain graph* (TSCG) harbors the conditional (in)dependencies.

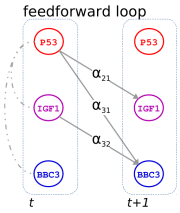

Temporal conditional independencies:

$Y_{j',t} \perp\!\!\!\perp Y_{j'',t+1} \,|\, \text{other } Y_{j''',t}\text{'s.}$

Parametric criterion:

$$[(\mathbf{A})^{-1}]_{j',j''} = 0 \iff Y_{j'',t} \perp\!\!\!\perp Y_{j',t+1} \,|\, \text{other } Y_{j''',t}\text{'s.}$$

Examples:

$Y_{\text{P53},t+1} \perp\!\!\!\perp Y_{\text{BBC3},t} \mid Y_{\text{IGF1},t}$ (*no* edge).
$Y_{\text{P53},t} \not\perp\!\!\!\perp Y_{\text{IGF1},t+1} \mid Y_{\text{BBC3},t}$ (*directed* edge).

Auto-regression

The *time-series chain graph* (TSCG) harbors the conditional (in)dependencies.

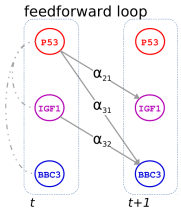

Contemporaneous cond. independencies:

$Y_{j',t+1} \perp\!\!\!\perp Y_{j'',t+1} \,|\, \text{other } Y_{j''',t+1}\text{'s, all } Y_{j,t}\text{'s.}$

Parametric criterion:

$$(\mathbf{\Omega}_\varepsilon)_{j',j''} = 0 \iff Y_{j',t+1} \perp\!\!\!\perp Y_{j'',t+1} \,|\, \text{other } Y_{j''',t+1}\text{'s, all } Y_{j,t}\text{'s.}$$

Examples:

$Y_{\text{IGF1},t+1} \perp\!\!\!\perp Y_{\text{BBC3},t+1} \mid Y_{\text{TP53},t+1}, Y_{\text{IGF1},t}, Y_{\text{BBC3},t}, Y_{\text{TP53},t}$ (*no* edge).
$Y_{\text{P53},t+1} \not\perp\!\!\!\perp Y_{\text{IGF1},t+1} \mid Y_{\text{BBC3},t+1}, Y_{\text{IGF1},t}, Y_{\text{BBC3},t}, Y_{\text{TP53},t}$ (*undirected* edge).

Auto-regression

Global conditional (in)dependencies:

$Y_{j',*} \perp\!\!\!\perp Y_{j'',*} \,|\, \text{other } Y_{j''',*}\text{'s}$

Wermuth condition for global conditional independency:
- unconnected nodes do not exert influence on the same node.

Consider a four-gene VAR(1) process. Then, global conditional independence of genes 1 and 2 forbids motifs like:

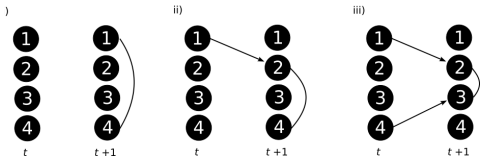

Goals

- learn VAR(1) model,
- reconstruct TSCG,
- exploit VAR(1) model and TSCG, and
- explore extensions.

Data (Section 1.2)

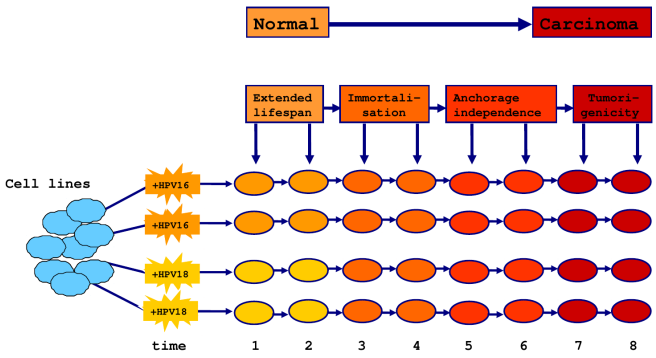

Data (Section 1.2)

The *central dogma of molecular biology* discerns three steps in the process that leads from the information encoded in DNA to the proteins in the cell:
- Replication
- Transcription
- Translation

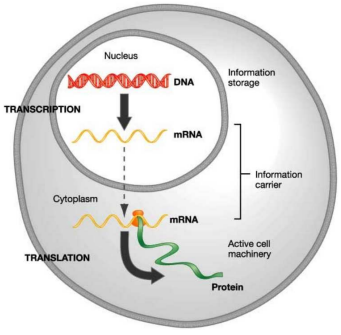

Gene expression data
is proportional to # mRNA molecules present in the cell.

Data (Section 1.2)

Load gene expression data related to P53 signalling pathway.

```
# load data
data(hpvP53)

# reformat data
Y <- longitudinal2array(t(exprs(hpvP53rna)))

# inquire design
pData(hpvP53rna)
fData(hpvP53rna)
```

The `Y` is an array-class object.

```
# interrogate dimension
> dim(Y)
```

Its dimensions: # genes × # time points × # cell lines

Data (Section 1.3)

Visualize data of single gene.

```
# plot time-courses of a single gene
plotVAR1data(Y[5, , , drop=FALSE])
```

Visualize data of single cell line.

```
# K-means clustering of the genes
cellLine <- 1
kClust   <- kmeans(Y[, , cellLine],
                   centers=7, nstart=100)$cluster

# heatmap of reshuffled data
edgeHeat(Y[unlist(lapply(1:max(kClust),
                  function(id, clusters){
                  which(clusters==id) },
                  kClust)), , cellLine])
```

The VAR(1) model (Section 1.3)

Q: the study comprises how many genes, cell lines, and time points?

Q: the BBC3 gene lies on which chromosome?

Q: slicing, e.g., `Y[ , , 1]`, an array objects yields ...?

Q: are there genes that are upregulated over time?

Q: are there genes that exhibit different behaviour between cell lines?

Q: are there groups of genes that exhibit similar behaviour over time?

Data (Section 1.3, answer)

Experimental design

```
    cellLine hpv time
1-1        1  16    1
1-2        2  16    1
1-3        3  18    1
1-4        4  18    1
2-1        1  16    2
2-2        2  16    2
...
```

Annotation information

```
                gene entrezID         probe chr
APAF1          APAF1      317  A_23_P36611  12
ATM              ATM      472  A_23_P35916  11
ATR              ATR      545 A_23_P136058   3
BAX              BAX      581 A_23_P208706  19
BBC3            BBC3    27113 A_24_P305312  19
BID              BID      637 A_23_P154929  22
...
```

Data (Section 1.3, answer)

First impression of the data

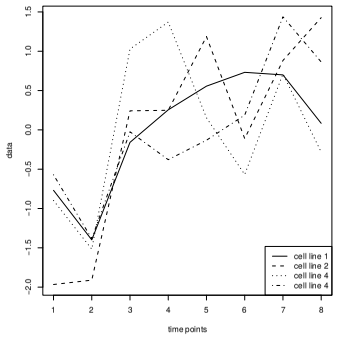

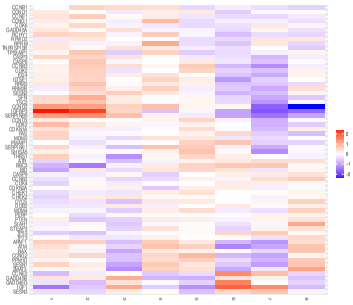

Estimation

Experimental design

- $j = 1, \ldots, p$ genes,
- $t = 1, \ldots, \mathcal{T}$ time points,
- $i = 1, \ldots, n$ cell lines.

Here, $p = 64$, $\mathcal{T} = 8$, and $n = 4$.

Data

- $Y_{j,t,i}$ is expression of gene j at time point t of cell line i.

Assumptions

- cell lines are independent,
- cell line data $\sim$ same VAR(1) model.

Estimation

The likelihood of cell line i:

$$P(\mathbf{Y}_{*,\mathcal{T},i}, \mathbf{Y}_{*,\mathcal{T}-1,i}, \ldots, \mathbf{Y}_{*,1,i}) \quad = \quad P(\mathbf{Y}_{*,1,i}) \prod_{t=1}^{\mathcal{T}-1} P(\mathbf{Y}_{*,t+1,i} \,|\, \mathbf{Y}_{*,t,i}),$$

due to the 1st order Markov assumption.

Furthermore, due to the VAR(1) model assumption,

$$P(\mathbf{Y}_{*,t+1,i} \,|\, \mathbf{Y}_{*,t,i}) \quad = \quad \phi_{\mathbf{A}\mathbf{Y}_{*,t,i}, \mathbf{\Omega}_\varepsilon^{-1}}(\mathbf{Y}_{*,t+1,i}).$$

Together with the independence assumption, this yields the likelihood:

$$\propto \prod_{i=1}^{n} \prod_{t=1}^{\mathcal{T}-1} (2\pi)^{-p/2} |\mathbf{\Omega}_\varepsilon|^{1/2} \exp[-\tfrac{1}{2}(\mathbf{Y}_{*,t+1,i} - \mathbf{A}\mathbf{Y}_{*,t,i})^\top \mathbf{\Omega}_\varepsilon (\mathbf{Y}_{*,t+1,i} - \mathbf{A}\mathbf{Y}_{*,t,i})].$$

Q: why can we drop the 'start probability'?

Estimation

The ridge ML estimator of the VAR(1) model maximizes

$$\text{loglikelihood} - \lambda_a \|\text{vec}(\mathbf{A})\|_2^2 - \lambda_\omega \|\mathbf{\Omega}_\varepsilon\|_2^2.$$

Targets $\mathbf{A}_0$ and $\mathbf{\Omega}_0$ could be included.

Analytic expressions of estimators are available:

$$\begin{aligned} \widehat{\mathbf{A}}(\lambda_a) \quad &= \quad [\lambda_a \mathbf{I}_{p^2 p^2} + \widehat{\mathbf{\Gamma}}(0) \otimes \mathbf{\Omega}_\varepsilon]^{-1} \text{vec}[\mathbf{\Omega}_\varepsilon \widehat{\mathbf{\Gamma}}(-1)], \\ \widehat{\mathbf{\Omega}}_\varepsilon(\lambda_\omega) \quad &= \quad [\tfrac{1}{2}\mathbf{S}_\varepsilon + (\lambda_\omega \mathbf{I}_{pp} + \tfrac{1}{4}\mathbf{S}_\varepsilon^2)^{1/2}]^{-1}, \end{aligned}$$

with sufficient statistics $\widehat{\mathbf{\Gamma}}(0)$ and $\widehat{\mathbf{\Gamma}}(-1)$ and

$$\mathbf{S}_\varepsilon \quad = \quad n^{-1}(\mathcal{T}-1)^{-1} \sum_{i=1}^{n} \sum_{t=1}^{\mathcal{T}-1} (\mathbf{Y}_{*,t+1,i} - \mathbf{A}\mathbf{Y}_{*,t,i})(\mathbf{Y}_{*,t+1,i} - \mathbf{A}\mathbf{Y}_{*,t,i})^\top.$$

Alternate between updating of $\widehat{\mathbf{A}}(\lambda_a)$ and $\widehat{\mathbf{\Omega}}_\varepsilon(\lambda_\omega)$.

Estimation

The ridge ML estimators of $\mathbf{A}$ and $\mathbf{\Omega}_\varepsilon$

- are consistent in the $n \to \infty$-limits.
  If $\lambda_a, \lambda_\omega \xrightarrow{P} 0$, then $\widehat{\mathbf{A}}(\lambda_a) \xrightarrow{P} \mathbf{A}$ and $\widehat{\mathbf{\Omega}}_\varepsilon(\lambda_\omega) \xrightarrow{P} \mathbf{\Omega}_\varepsilon$.
- have Bayesian counterparts.
  E.g, if $\text{vec}(\mathbf{A}) \,|\, \mathbf{\Omega}_\varepsilon \sim \mathcal{N}(\mathbf{0}_{p^2}, \lambda_a^{-1} \mathbf{I}_{p^2 p^2})$, then $\mathbb{E}(\mathbf{A} \,|\, \mathbf{Y}, \mathbf{\Omega}_\varepsilon) = \widehat{\mathbf{A}}(\lambda_a)$.
- can be obtained by data augmentation.

Estimation

The auto-regression parameter $\mathbf{A}$ can also be estimated by means of ridge penalized sum-of-squares minimization:

$$\sum_{i=1}^{n} \sum_{t=1}^{\mathcal{T}-1} \|\mathbf{Y}_{*,t+1,i} - \mathbf{A}\mathbf{Y}_{*,t,i}\|_2^2 - \lambda_a \|\text{vec}(\mathbf{A})\|_2^2.$$

A target $\mathbf{A}_0$ could be included.

An analytic expressions of the estimator is available:

$$\begin{aligned} \widehat{\mathbf{A}}^{(\text{ridge ss})}(\lambda_a) \quad &= \quad [\lambda_a \mathbf{I}_{p^2 p^2} + \widehat{\mathbf{\Gamma}}(0) \otimes \mathbf{I}_{pp}]^{-1} \text{vec}[\widehat{\mathbf{\Gamma}}(-1)], \\ &\quad \text{vs.} \\ \widehat{\mathbf{A}}^{(\text{ridge ml})}(\lambda_a) \quad &= \quad [\lambda_a \mathbf{I}_{p^2 p^2} + \widehat{\mathbf{\Gamma}}(0) \otimes \mathbf{\Omega}_\varepsilon]^{-1} \text{vec}[\mathbf{\Omega}_\varepsilon \widehat{\mathbf{\Gamma}}(-1)]. \end{aligned}$$

Q: what is the advantage of the penalized sum-of-squares estimator?

Estimation (Section 1.5)

Code of ridge ML estimation

```
# fit the model
VAR1hat <- ridgeVAR1(Y=Y,
                     lambdaA=optLambdas[1],
                     lambdaP=optLambdas[2],
                     fitA   ="ml",
                     targetA=matrix(0, dim(Y)[1], dim(Y)[1]),
                     targetP=matrix(0, dim(Y)[1], dim(Y)[1]),
                     targetPtype="none")

# extract parameter estimates
Ahat <- VAR1hat$A
Phat <- VAR1hat$P

# visually inspect the estimates, e.g.
edgeHeat(Ahat, main="ridge estimate of A")
```

Estimation (Section 1.4)

Cross-validation chooses λ to optimize prediction.

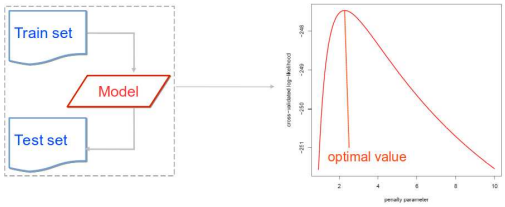

- Leave-one-out cross-validation leaves out time points one-by-one.
- Model learnt on training set.
- Fitted model applied to the left-out time point and evaluate the error.
- Average of these errors estimates the 'full data' predictor's error rate.

Estimation (Section 1.4)

R-code of cross-validation

```
# find optimal penalty parameters
optLambdas <- optPenaltyVAR1(Y,
                        lambdaMin =c(0.01, 0.00001),
                        lambdaMax =c(1000, 1),
                        lambdaInit=c(100 , 0.1),
                        optimizer ="nlm", ...)
```

Estimation (Section 1.4)

Investigate optimality of λ choice

```
# define grid for both penalty parameters
lambdaAgrid <- seq(300,    550, length.out=20)
lambdaPgrid <- seq(0.001, 0.01, length.out=20)

# evaluate LOOCV log-likelihood over the grid
LOOCVres <- loglikLOOCVcontourVAR1(lambdaAgrid,
                                   lambdaPgrid,
                                   Y, ...)
```

Plot the LOOCV log-likelihood contour with optimal penalty parameters:

```
# plot LOOCV log-likelihood somewhat nicer
contour(lambdaAgrid, lambdaPgrid, LOOCVres$llLOOCV, ...)

# add optimal penalty parameters
points(optLambdas[1], optLambdas[2], ...)
```

The VAR(1) model (Section 1.4 & 1.5)

Q: tune the penalty parameters.

Q: estimate the VAR(1) model parameters.

Q: investigate the optimality of the chosen penalty parameters.

Estimation (Section 1.4, answer)

LOOCV: $\lambda_a^{(\text{opt})} = 427.95465$ and $\lambda_\omega^{(\text{opt})} = 0.00259$.

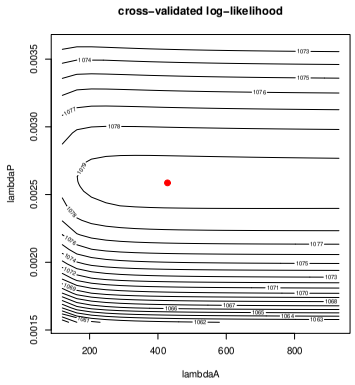

Estimation (Section 1.5, answer)

Heatmap of ridge ML estimates of $\mathbf{A}$ (left) and $\mathbf{\Omega}_\varepsilon$ (right).

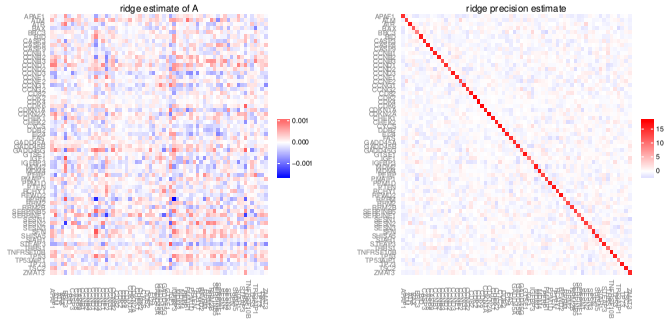

Estimation (Section 1.6)

Assess stability of $\widehat{\mathbf{A}}(\lambda_a)$.

- Perturb data: remove a design point.
- Re-do estimation, including LOOCV.
- Store elements of $\widehat{\mathbf{A}}_{-(t,i)}(\lambda_a)$.
- Repeat for all design points.
- Plot $\widehat{\mathbf{A}}_{-(t,i)}(\lambda_a)$'s vs $\widehat{\mathbf{A}}(\lambda_a)$.

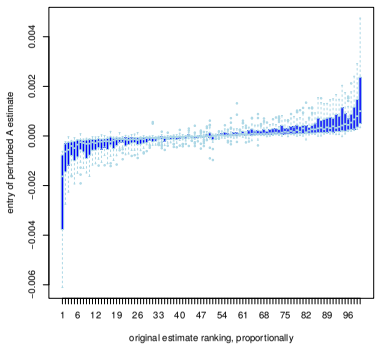

Estimation (Section 1.7)

Support determination

standardized Ω_ε estimate

	1	2	3	4	5
1	1.00	0.51	0.42	-0.03	0.34
2	*	1.00	0.40	0.11	-0.09
3	*	*	1.00	-0.57	0.04
4	*	*	*	1.00	0.07
5	*	*	*	*	1.00

model *as mixture*

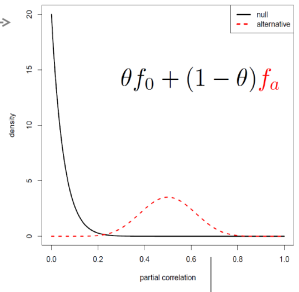

inferred Ω_ε support

	1	2	3	4	5
1	1	1	1	0	1
2	*	1	1	0	0
3	*	*	1	1	0
4	*	*	*	1	0
5	*	*	*	*	1

cut-off $\leftarrow P\left[(1,2) \notin \mathcal{E}_{\mathcal{G}} \,|\, (\widehat{\mathbf{\Omega}})_{1,2}\right]$

Estimation (Section 1.7)

R-code of sparsification

```
# support determination of A
zerosA <- sparsifyVAR1(A          =Ahat,
                       SigmaE     =symm(solve(Phat)),
                       threshold  ="top",
                       top        =25,
                       statistics =FALSE,
                       verbose    =FALSE)$zeros

# support determination of precision matrix
zerosP <- sparsify(P         =Phat,
                   threshold ="top",
                   top       =10,
                   output    ="light",
                   verbose   =FALSE)$zeros
```

Estimation (Section 1.8)

We re-estimate the parameters with known support.

This undoes some of the shrinkage.

The ridge ML estimates with known support maximize

$$\arg \max_{\substack{\{\mathbf{A}\,:\,\mathsf{C}_a\,\mathrm{vec}(\mathbf{A})=\mathbf{0}_{q_a}\} \\ \{\mathbf{\Omega}_\varepsilon\,:\,\mathsf{C}_\omega\,\mathrm{vec}(\mathbf{\Omega}_\varepsilon)=\mathbf{0}_{q_\omega}\}}} \text{loglikelihood} - \lambda_a \|\mathrm{vec}(\mathbf{A})\|_2^2 - \lambda_\omega \|\mathbf{\Omega}_\varepsilon\|_2^2.$$

The penalty parameters λ_a and λ_ω are re-tuned.

Estimation

R-code for calibration, re-tuning of the λ's.

```
# format precision support
supportP <- support4ridgeP(zeros =zerosP,
                           nNodes=nrow(Y))

# optimal penalty parameter determination
optLambdas <- optPenaltyVAR1(Y,
                             lambdaMin  =c(10^(-5), 10^(-5)),
                             lambdaMax  =c(10,       0.1),
                             lambdaInit =c(5,        0.01),
                             zerosA     =zerosA,
                             zerosP     =zerosP,
                             cliquesP   =supportP$cliques,
                             separatorsP=supportP$separators,
                             zerosAfit  ="sparse")
```

Estimation

R-code for calibration, re-estimation.

```
# re-fit the model
VAR1hat <- ridgeVAR1(Y          =Y,
                     lambdaA    =optLambdas[1],
                     lambdaP    =optLambdas[2],
                     zerosA     =zerosA,
                     cliquesP   =supportP$cliques,
                     separatorsP=supportP$separators,
                     zerosP     =zerosP,
                     zerosAfit  ="sparse")

# extract parameter estimates
Ahat <- VAR1hat$A
Phat <- VAR1hat$P

# heatmap of support of A
edgeHeat(adjacentMat(Ahat), ...)
```

Estimation (Section 1.8)

R-code to plot the reconstructed time-series chain graph.

```
# time-series chain graph
graphVAR1(Ahat,
          Phat,
          nNames=rownames(Ahat),
          type  ="TSCG",
          ...)
```

R-code to plot the (global) conditional independence graph.

```
# determine adjacency matrix
adjMatCIG <- CIGofVAR1(Ahat, Phat)

# plot global conditional independence graph
graphVAR1(Ahat,
          Phat,
          nNames=rownames(Ahat),
          type  ="globalPC")
```

The VAR(1) model (Section 1.7 & 1.8)

Q: sparsify the estimated VAR(1) model parameters.

Q: re-tune the penalty parameters with inferred support.

Q: re-estimate the VAR(1) model parameters with inferred support.

Q: visualize model parameter.

Q: draw the time-series chain graph.

The VAR(1) model (Section 1.7, answer)

For the autoregression parameter we obtain

```
Step 1... determine cutoff point
Step 2... estimate parameters of null distribution and eta0
Step 3... compute p-values and estimate empirical PDF/CDF
Step 4... compute q-values and local fdr
Step 5... prepare for plotting

-> Retained elements:  134
-> Corresponding to 3.27 % of possible elements
```

while for the error precision matrix

```
...

- Retained elements:  14
- Corresponding to 0.69 % of possible edges
```

Estimation (Section 1.8, answer)

LOOCV: $\lambda_a^{(\text{opt})} = 2.55082$ and $\lambda_\omega^{(\text{opt})} = 0.00583$, known support.

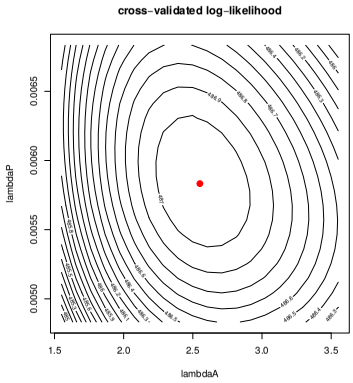

Estimation (Section 1.8, answer)

Heatmap of ridge ML estimates of $\mathbf{A}$ (left) and $\mathbf{\Omega}_\varepsilon$ (right), known support.

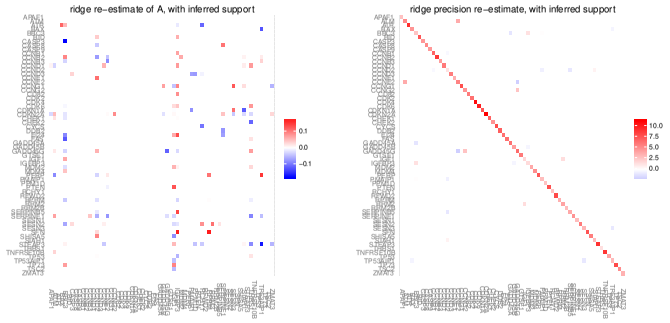

Estimation (Section 1.8, answer)

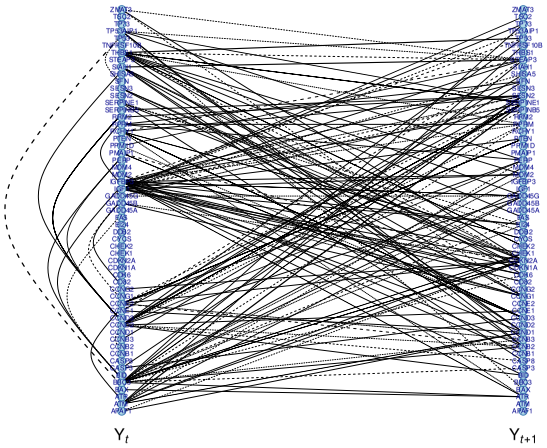

Fit (Section 1.9)

The fit is

$$\widehat{\mathbf{Y}}_{*,t+1,i} \quad = \quad \widehat{\mathbf{A}}(\lambda_a)\mathbf{Y}_{*,t,i}.$$

R-code to obtain the fit.

```
# calculate fits and fit-vs-observation correlations
Yhat <- array(dim=dim(Y))
for (i in 1:dim(Y)[3]){
  Yhat[, -1, i] <- Ahat %*% Y[, -dim(Y)[2], i]
}
```

Ideally, $\widehat{\mathbf{Y}}_{*,t,i}$ is close (in some sense) to $\mathbf{Y}_{*,t,i}$.

Fit (Section 1.9)

R-code to compare $\widehat{\mathbf{Y}}_{j,*,i}$ and $\mathbf{Y}_{j,*,i}$ visually.

```
# specify molecular entity of interest
entityID <- 1

# format data for plotting
label <- paste("Cell", sort(rep(1:dim(Y)[3], dim(Y)[2]-1)))
...
```

R-code to compare $\mathrm{Cor}(\widehat{\mathbf{Y}}_{j,*,i}, \mathbf{Y}_{j,*,i})$.

```
# calculate fits and fit-vs-obs correlations
corFit <- numeric()
for (j in 1:dim(Y)[1]){
  ...
}

# histogram of the fit-vs-obs correlations
hist(corFit, ...)
```

Fit (Section 1.9)

Q: calculate the fit and compare it to the observations.

Q: why is gene 1 a 'flatliner'?

Q: how does it compare to the other histograms (correlation with previous observation and overall mean)?

Fit (Section 1.9, answer)

Scatter plots of $\widehat{\mathbf{Y}}_{j,*,i}$ and $\mathbf{Y}_{j,*,i}$.

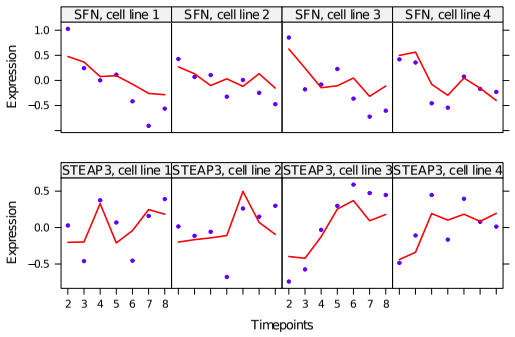

Fit (Section 1.9, answer)

Histogram of $\mathrm{Cor}(\widehat{\mathbf{Y}}_{j,*,i}, \mathbf{Y}_{j,*,i})$.

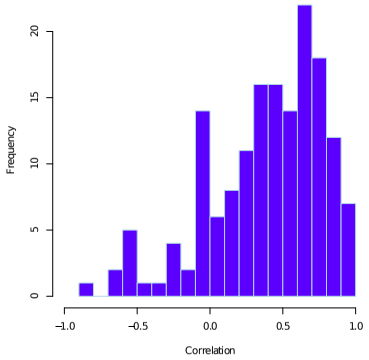

Exploitation (Section 1.10.1)

Calculate node-wise network statistics.

Take into account directionality.

For instance, in- and out-degree of a node.

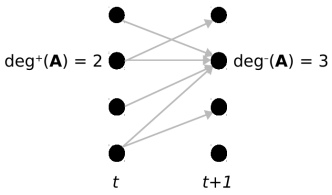

Code to calculate these statistics.

```
# calculate node-wise network stats
nodeStats <- nodeStatsVAR1(Ahat, Phat, as.table=TRUE)
```

Network statistics are based on both the TSCG and the global CIG graph!

Exploitation (Section 1.10.2)

Mutual information quantifies the information gain on one random variable by observing another. Formally,

$$\begin{aligned}\mathcal{I}(\mathbf{Y}_{*,t+\tau,i}, Y_{j,t,i} \,|\, \mathbf{Y}_{*,t-1,i}) &= \mathcal{H}(\mathbf{Y}_{*,t+\tau,i} \,|\, \mathbf{Y}_{*,t-1,i}) \\ &\quad -\mathcal{H}(\mathbf{Y}_{*,t+\tau,i} \,|\, Y_{j,t,i}, \mathbf{Y}_{*,t-1,i}),\end{aligned}$$

where $\mathcal{H}$ denotes the entropy.

Here the entropy is e.g.

$$\mathcal{H}(\mathbf{Y}_{*,t+\tau,i} \,|\, \mathbf{Y}_{*,t-1,i}) = \log\left\{\det\left[\mathrm{Var}(\mathbf{Y}_{*,t+\tau,i} \,|\, \mathbf{Y}_{*,t-1,i})\right]\right\}.$$

Code to plot the (global) conditional independence graph.

```
# evaluate mutual informations with specified lag
MIs <- mutualInfoVAR1(Ahat, solve(Phat), lag=1)
```

Exploitation (Section 1.10.3)

Impulse response analysis

Assess the change in variates at future time points due to an exogeneous change at the current time:

$$\frac{\partial \mathbf{Y}_{*,t+\tau}}{\partial \boldsymbol{\varepsilon}_{*,t}} = \mathbf{A}^{\tau}.$$

Facilitates prediction of a knockout effect.

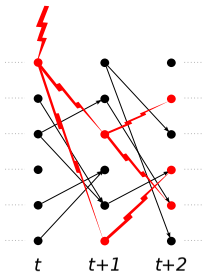

R-code

```
# evaluate impulse response with specified lag
IRs <- impulseResponseVAR1(Ahat, lag=1)
```

Exploitation (Section 1.10.4)

R-code to generate motif frequency table.

```
# motif detection
motifs <- motifStatsVAR1(Ahat)
```

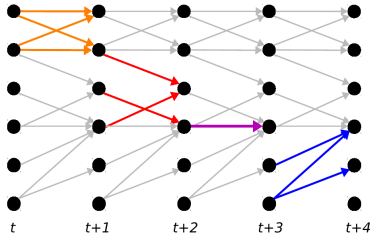

Auto-regulator	: 11
Coupled auto-regulators	: 3
Feedback loop	: 7
Feedforward loop	: 2
...	: ...

Exploitation (Section 1.10.5)

Covariance decomposition

Given TSCG, decompose conditional covariance in terms of paths as:

$$\mathrm{Cov}(Y_{j',t,i}, Y_{j'',t+\tau,i} \,|\, \mathbf{Y}_{*,t-1,i}) = (\boldsymbol{\Sigma}_{\varepsilon}\mathbf{A}^{\tau})_{j',j''} = \sum_{j'''=1}^{p} (\boldsymbol{\Sigma}_{\varepsilon})_{j',j'''} (\mathbf{A}^{\tau})_{j''',j''}.$$

E.g.

$$\mathrm{Cov}(\mathtt{BBC3}_t, \mathtt{TP73}_t \,|\, \ldots) = -0.002483485 - 0.0006845158$$

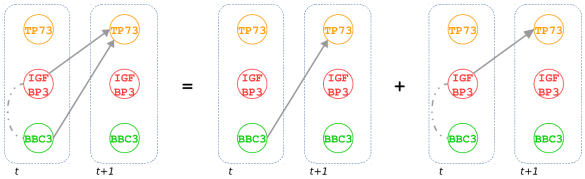

Fit (Section 1.10)

Q: Calculate the node-statistics.
On the basis of these statistics can you discern types of genes?

Q: The GADD45G gene is a potential tumor suppressor.
If activated, it may form a novel treatment.
Which other gene's perturbation at time t stimulates GADD45G most a $t+1$? And at $t+2$?

Q: What is the most prevalent motif?
Is there an *incoherent* feedforward loop present in the TSCG?

Exploitation (Section 1.10.1, answer)

Node-wise network statistics

	$\deg^-(\mathbf{A})$	$\deg^+(\mathbf{A})$	between.	close.	eigen centr.
BBC3	0	17	17	0.00192	1.00000
CCND2	0	12	18	0.00187	0.68783
IGF1	1	14	0	0.00191	0.97635
IGFBP3	0	16	7	0.00190	0.87513
THBS1	0	11	0	0.00188	0.87717
CCNG1	6	0	0	0.00177	0.25154
CDKN2A	12	0	0	0.00181	0.49508
SERPINE1	8	4	0	0.00185	0.70869
SESN2	8	0	0	0.00180	0.26759
STEAP3	9	0	0	0.00179	0.36285

upper five ≈ 'regulators'
lower five ≈ 'regulatees'

'regulator'

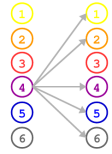

'regulatee'

Exploitation (Sections 1.10.2&3, answer)

Mutual information results
- 'regulator' ($\deg^+ = 17$) : mutual info of $\mathtt{BBC3}_t$ on $t+1 = 0.05605$,
- 'regulatee' ($\deg^- = 12$) : mutual info of $\mathtt{CDKN2A}_t$ on $t+1 = 0.00000$.

Result of impulse response analyses, e.g.:
- 'regulator' ($\deg^+ = 17$) : effect of $\mathtt{BBC3}_t$ on all $t+1 = 0.01497^*$,
- 'regulatee' ($\deg^- = 12$) : effect of $\mathtt{CDKN2A}_t$ on all $t+1 = 0.00000^*$.

* the average (over all genes) of their absolute changes due to a change in, e.g., BBC3 at the previous time point.

Exploitation (Sections 1.10.3, answer)

Study to impulse response of GADD45G due to a change in all others

```
id <- which(impulseResponseVAR1(Ahat, 1)[34,] != 0)
impulseResponseVAR1(Ahat, 1)[34, id]
```

At $t+1$:

```
    ATM     BBC3    CCND2 GADD45G   IGFBP3    THBS1
 0.0165 -0.0314   0.0563 -0.1136 -0.0278   0.0714
```

At $t+2$:

```
    ATM     BBC3    CCND2 GADD45G   IGFBP3    THBS1
-0.0019   0.0036 -0.0064   0.0129   0.0032 -0.0081
```

Exploitation (Section 1.10.4, answer)

The motif frequency table is plotted ...

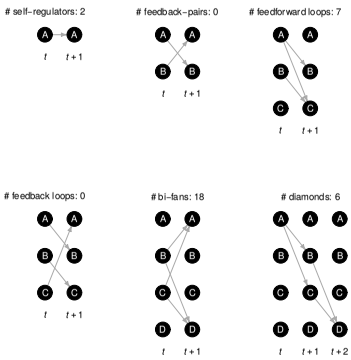

Exploitation (Section 1.10.4, answer)

... but also printed on screen:

```
motif summary table:
                motif frequency
 -> selfregulator     :         2
 -> feedback pair     :         0
 -> feedforward loop  :         7
 -> feedback loop     :         0
 -> bifan             :        18
 -> diamond           :         6
```

A `list` with details of all motifs is also provided. E.g.,

```
$feedforwardloops[[1]]
       t t+1 sign
path1  2   50    1
path2  2   25    1
path3 25   50    0
```

Recall: a motif's (in)coherence is determined by its signs.

Exploitation (Section 1.10.4, answer)

Motif analysis

E.g. feedforward motif

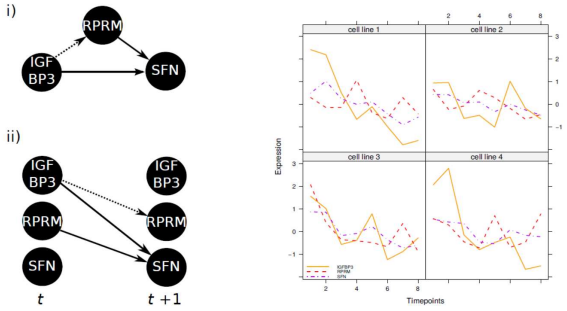

Extensions

Multiple VAR(1) models (Section 2)

Investigate difference between HPV16 and HPV18 transfected cell lines.

```
table(pData(hpvP53rna)[,1:2])
```

Both HPV types lead to cervical cancer. The regulatory networks of their transfected cell lines may share some similarity.

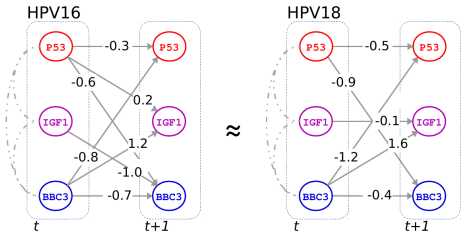

Multiple VAR(1) models (Section 2)

Assume a VAR(1) model per HPV-type, e.g.

$$\mathbf{Y}_{*,t+1,i} \quad = \quad \mathbf{A}_{\text{hpv16}}\mathbf{Y}_{*,t+1,i} + \boldsymbol{\varepsilon}_{*,t+1,i},$$

for i any of the HPV16 transfected cell line and $\boldsymbol{\varepsilon}_{*,t,i} \sim \mathcal{N}(\mathbf{0}_p, \boldsymbol{\Omega}_\varepsilon^{-1})$.

Thus, a *different* auto-regression parameter but *common* error.

The fused ridge ML estimator of the VAR(1) model maximizes

$$\begin{aligned}\text{loglikelihood} &- \lambda_a \|\text{vec}(\mathbf{A}_{\text{hpv16}})\|_2^2 - \lambda_a \|\text{vec}(\mathbf{A}_{\text{hpv18}})\|_2^2 \\ &- \lambda_f \|\text{vec}(\mathbf{A}_{\text{hpv16}}) - \text{vec}(\mathbf{A}_{\text{hpv18}})\|_2^2 - \lambda_\omega \|\boldsymbol{\Omega}_\varepsilon\|_2^2.\end{aligned}$$

Targets $\mathbf{A}_0$ (common to both HPV types) and $\boldsymbol{\Omega}_0$ could be included.

Multiple VAR(1) models (Section 2)

R-code for fused ridge estimation:

```
# group indicies
id <- c(rep(1, 2), rep(2, 2)) - 1

# search for optimal penalty parameters
optLambdas <- optPenaltyVAR1fused(Y=Y, id=id, ...)

# fit fused VAR(1) model
VAR1hatF <- ridgeVAR1fused(Y=Y, id=id, ...)
Ahat16   <- VAR1hatF$As[1:64,]
```

After support inferral of $\mathbf{A}_{\text{hpv16}}$ and $\mathbf{A}_{\text{hpv18}}$, continue either by

- estimation per HPV-type, or
- estimation jointly.

Down-stream analysis proceeds per HPV-type.

Multiple VAR(1) models (Section 2)

Q: jointly estimate the VAR(1) models by means fused ridge estimation.
Spoiler: $(\lambda_a^{\text{opt}}, \lambda_f^{\text{opt}}, \lambda_\omega^{\text{opt}}) = (19.4444, 21.6667, 0.0024)$.

Q: sparsify and callibrate the VAR(1) models.

Q: draw heatmaps of the VAR(1) model parameters.

Q: draw the time-series chain graphs of the VAR(1) models.

Q: investigate differences in regulation structure.

Multiple VAR(1) models (Section 2, answer)

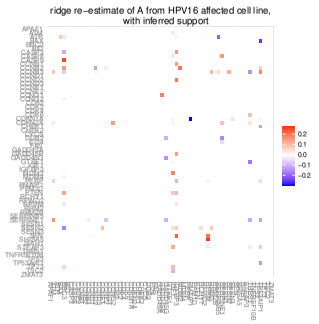

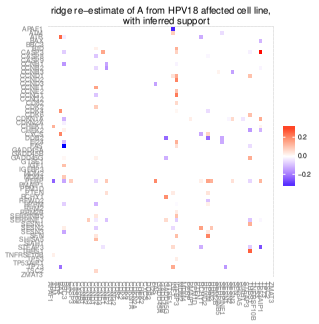

Multiple VAR(1) models (Section 2, answer)

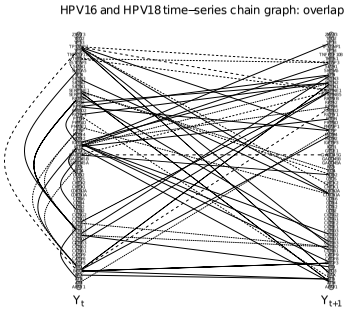

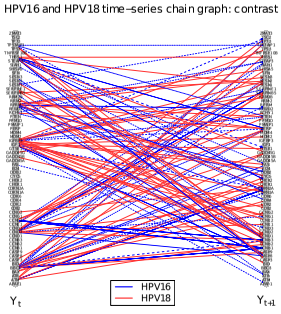

$\#\{(j,j') : (\widehat{\mathbf{A}}_{\text{hpv16}})_{j,j'} \neq 0\} = 88$, and
$\#\{(j,j') : (\widehat{\mathbf{A}}_{\text{hpv18}})_{j,j'} \neq 0\} = 125$

Multiple VAR(1) models (Section 2, answer)

Regulators $\approx$ the same, while regulatees (e.g. PERP) exhibit some differences.

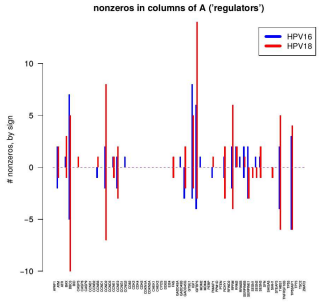

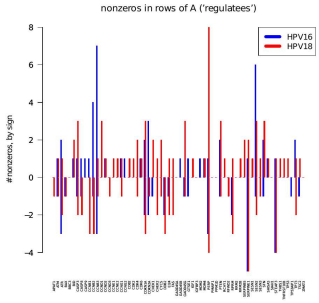

Q: is this due to differences in # selected elements of $\mathbf{A}$?

The VAR(2) model (Section 3)

The cell line data may be better modelled by a process with a longer memory.

A VAR(2) process:

$$\begin{aligned} Y_{1,t+1} &= \varepsilon_{1,t+1}, \\ Y_{2,t+1} &= \alpha_{21}^{(1)} Y_{1,t} + \alpha_{13}^{(2)} Y_{3,t-1} + \varepsilon_{2,t+1}, \\ Y_{3,t+1} &= \alpha_{31}^{(1)} Y_{1,t} + \alpha_{32}^{(1)} Y_{2,t} + \varepsilon_{3,t+1}. \end{aligned}$$

In matrix notation

$$\mathbf{Y}_{t+1} = \mathbf{A}_1 \mathbf{Y}_t + \mathbf{A}_2 \mathbf{Y}_{t-1} + \varepsilon_{t+1}$$

with $\varepsilon_t \sim \mathcal{N}(\mathbf{0}_p, \mathbf{\Omega}_\varepsilon^{-1})$.

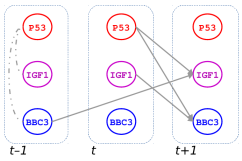

The VAR(2) model

The VAR(2) model can be reformulated as a VAR(1) model:

$$\begin{pmatrix} \mathbf{Y}_{t+1} \\ \mathbf{Y}_t \end{pmatrix} = \begin{pmatrix} \mathbf{A}_1 & \mathbf{A}_2 \\ \mathbf{I}_{pp} & \mathbf{0}_{pp} \end{pmatrix} \begin{pmatrix} \mathbf{Y}_t \\ \mathbf{Y}_{t-1} \end{pmatrix} + \begin{pmatrix} \varepsilon_{t+1} \\ \mathbf{0}_p \end{pmatrix}.$$

The same applies to VAR-type models with a higher lag.

The VAR(2) model (Section 3)

R-code:

```
# search for optimal penalty parameters
optLambda <- optPenaltyVAR2(Y, ...)

# fit VAR(2) model
VAR2hat <- ridgeVAR2(Y=Y, lambdaA1=optLambda[1],
                          lambdaA2=optLambda[2], ...)
A1hat <- VAR2hat$A1
A2hat <- VAR2hat$A2
...

# determine support for As and P
zeros <- sparsifyVAR2(Ahat1, Ahat2,
                      SigmaE=symm(solve(Phat)), ...)

# node statistics table
stats <- nodeStatsVAR2(Ahat1, Ahat2, Phat, ...)
```

The VAR(2) model (Section 3)

Q: estimate, sparsify, and callibrate the VAR(2) model.

Q: draw the time-series chain graph of the VAR(2) model.

Q: investigate whether the inclusion of the second lag term $\mathbf{A}_2 \mathbf{Y}_{t-1}$ is worthwhile.

The VAR(2) model (Section 3, answer)

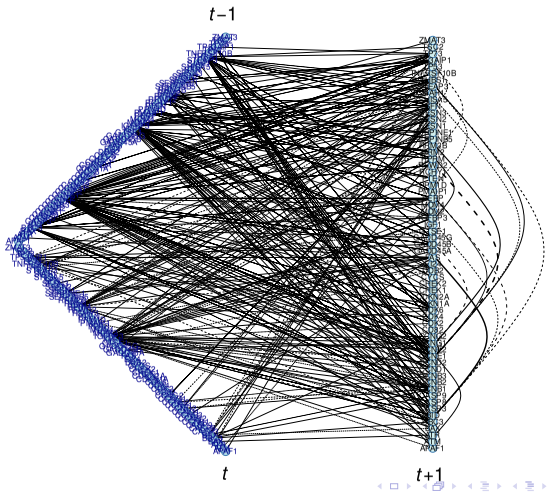

The VAR(2) model (Section 3, answer)

Much smaller elements of $\widehat{\mathbf{A}}_2$ than $\widehat{\mathbf{A}}_1$ as:

initial	$\lambda_{a_1} = 0.0398$	$\lambda_{a_2} = 13.0090$	$\lambda_\omega = 0.0068$
calibration	$\lambda_{a_1} = 0.3310$	$\lambda_{a_2} = 2.0324$	$\lambda_\omega = 0.0063$

The VAR(1) model with
$\#\{(j,j') : (\widehat{\mathbf{A}})_{j,j'} \neq 0\} = 134\}$.
is more parsimonious than
the VAR(2) model with
$\#\{(j,j') : (\widehat{\mathbf{A}}_1)_{j,j'} \neq 0\} = 107\}$ and
$\#\{(j,j') : (\widehat{\mathbf{A}}_2)_{j,j'} \neq 0\} = 138\}$.

The 'fit' of the VAR(1) and VAR(2) are comparable.

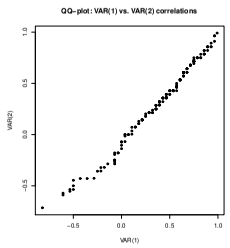

The VARX(1) model (Section 4)

Incorporate time-varying covariates into the VAR(1) model.

A VARX(1) process:

$$
\begin{aligned}
Y_{1,t+1} &= \beta_{11} X_{1,t+1} + \qquad\qquad \varepsilon_{1,t+1}, \\
Y_{2,t+1} &= \beta_{22} X_{2,t+1} + \alpha_{21} Y_{1,t} + \qquad \varepsilon_{2,t+1}, \\
Y_{3,t+1} &= \beta_{33} X_{3,t+1} + \alpha_{31} Y_{1,t} + \alpha_{32} Y_{2,t} + \varepsilon_{3,t+1}.
\end{aligned}
$$

In matrix notation

$$\mathbf{Y}_{t+1} = \mathbf{B}\mathbf{X}_{t+1} + \mathbf{A}\mathbf{Y}_t + \varepsilon_{t+1}$$

with $\varepsilon_t \sim \mathcal{N}(\mathbf{0}_p, \mathbf{\Omega}_\varepsilon^{-1})$.

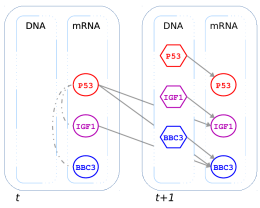

The VARX(1) model (Section 4)

DNA copy number : # copies of a genomic segment present in cell.

- 2 : autosomal chromosomes,
- 1 : X or Y chrom. in males,
- 0 : Y chromosome in females,
- anything goes in cancer.

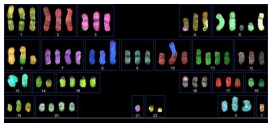

A proportionality law:
More genomic (=DNA) copies of a gene
$\Longrightarrow$ more product (=mRNA) it encodes for is formed.

Debate
Is proportionality a localized phenomenon?

The VARX(1) model (Section 4)

We allow only the *cis* (and not the *trans*) effect

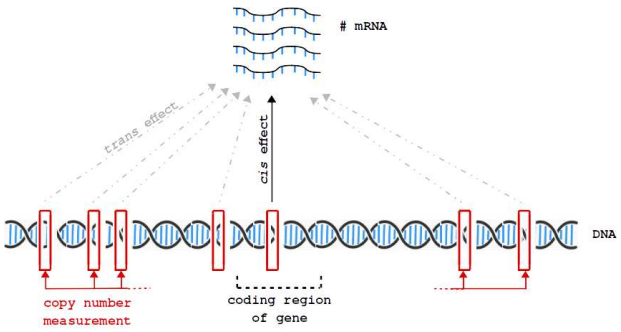

This implies a diagonal support of $\mathbf{B}$.

The VARX(1) model (Section 4)

R-code:

```
# load covariate data
X <- longitudinal2array(t(copynumber(hpvP53cn)))

# load the constrained elements to zero for B
zerosB <- which(cn2rna==0, arr.ind=TRUE)

# search for optimal penalty parameters
optLambda <- optPenaltyVARX1(Y=Y, X=X, lagX=0,
                             zerosB=zerosB, ...)

# fit VAR(1) model
VAR1hat <- ridgeVARX1(Y=Y, X=X,
                      lambdaA=optLambda1[1],
                      lambdaB=optLambda1[2],
                      lambdaP=optLambda1[3],
                      lagX=0, zerosB=zerosB)
Ahat <- VAR1hat$A
Bhat <- VAR1hat$B
...
```

The VARX(1) model (Section 4)

Q: estimate, sparsify, and callibrate the VARX(1) model.

Q: draw the time-series chain graph of the VARX(1) model.

Q: investigate the proportionality law.

Q: is there complementary regulation?

The VARX(1) model (Section 4, answers)

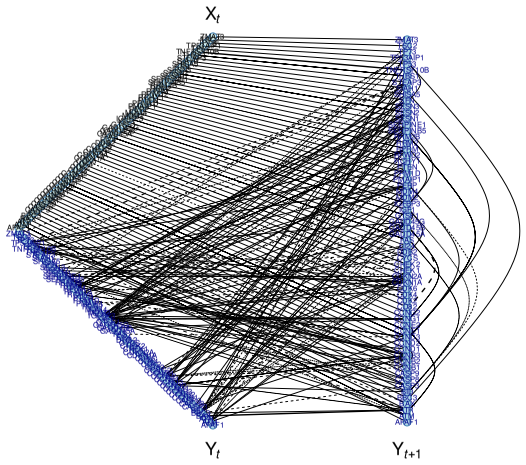

The VARX(1) model (Section 4, answers)

Corroboration of the proportionality law.

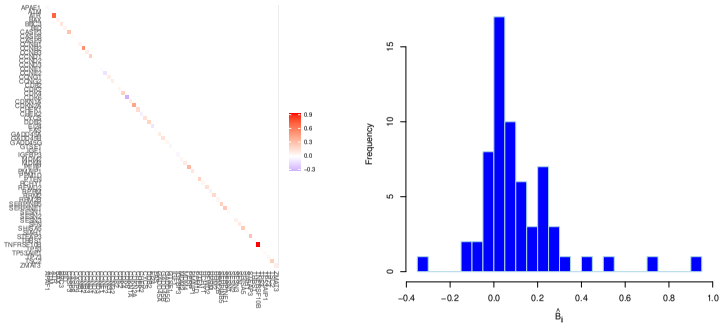

The VARX(1) model (Section 4, answers)

Complementary regulation mechanisms?

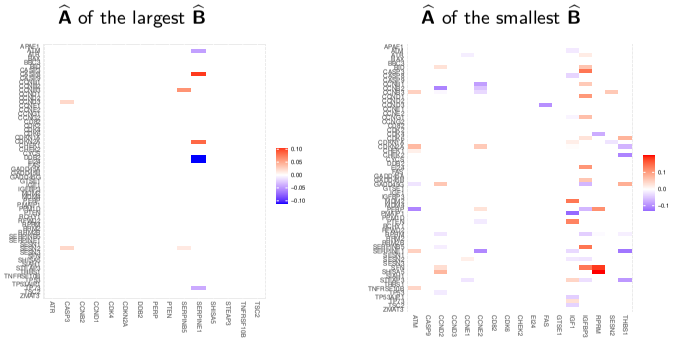

References

These slides are based on:

- Miok, V., Wilting, S.M., van Wieringen, W.N. (2017). Ridge estimation of the VAR(1) model and its time series chain graph from multivariate time-course omics data. *Biometrical Journal*, 59(1), 172–191
- Miok, V., Wilting, S.M., van Wieringen, W.N. (2019). Ridge estimation of network models from time-course omics data. *Biometrical Journal*, 61(2), 391–405.
- van Wieringen, W.N. (2020) ragt2ridges: Ridge Estimation of Vector Auto-Regressive (VAR) Processes. R package version 0.3.4. `https://CRAN.R-project.org/package=ragt2ridges`.

License